

# nhanes\_univariate\_practice

October 18, 2022

## 1 Practice notebook for univariate analysis using NHANES data

This notebook will give you the opportunity to perform some univariate analyses on your own using the NHANES. These analyses are similar to what was done in the week 2 NHANES case study notebook.

You can enter your code into the cells that say “enter your code here”, and you can type responses to the questions into the cells that say “Type Markdown and Latex”.

Note that most of the code that you will need to write below is very similar to code that appears in the case study notebook. You will need to edit code from that notebook in small ways to adapt it to the prompts below.

To get started, we will use the same module imports and read the data in the same way as we did in the case study:

```
In [1]: %matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import statsmodels.api as sm
import numpy as np

da = pd.read_csv("nhanes_2015_2016.csv")
```

```
In [4]: print(da.info())
data_size = da.shape
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5735 entries, 0 to 5734
Data columns (total 28 columns):
SEQN          5735 non-null int64
ALQ101        5208 non-null float64
ALQ110        1731 non-null float64
ALQ130        3379 non-null float64
SMQ020        5735 non-null int64
RIAGENDR      5735 non-null int64
RIDAGEYR      5735 non-null int64
RIDRETH1      5735 non-null int64
DMDCHITZN     5734 non-null float64
```

```

DMDDEDUC2    5474 non-null float64
DMDMARTL     5474 non-null float64
DMDHHSIZ     5735 non-null int64
WTINT2YR     5735 non-null float64
SDMVPSU      5735 non-null int64
SDMVSTRA     5735 non-null int64
INDFMPPIR    5134 non-null float64
BPXSY1       5401 non-null float64
BPXDI1       5401 non-null float64
BPXSY2       5535 non-null float64
BPXDI2       5535 non-null float64
BMXWT        5666 non-null float64
BMXHT        5673 non-null float64
BMXBMI       5662 non-null float64
BMXLEG       5345 non-null float64
BMXARML      5427 non-null float64
BMXARMC      5427 non-null float64
BMXWAIST     5368 non-null float64
HIQ210       4732 non-null float64
dtypes: float64(20), int64(8)
memory usage: 1.2 MB
None

```

```

In [2]: # Check the class distribution of DMDMARTL using groupby and size
        print(da.groupby('DMDMARTL').size())

```

```

DMDMARTL
1.0      2780
2.0       396
3.0       579
4.0       186
5.0      1004
6.0       527
77.0        2
dtype: int64

```

## 1.1 Question 1

Relabel the marital status variable `DMDMARTL` to have brief but informative character labels. Then construct a frequency table of these values for all people, then for women only, and for men only. Then construct these three frequency tables using only people whose age is between 30 and 40.

```

In [5]: # insert your code here
        da["DMDMARTLx"] = da.DMDMARTL.replace({1: "Married", 2: "Widowed", 3: "Divorced", 4: "S
        da.DMDMARTLx.value_counts()

```

```
Out [5]: Married                2780
        Never Married          1004
        Divorced                579
        Living with Partner     527
        Widowed                 396
        Missing                 261
        Separated               186
        Refused                  2
        Name: DMDMARTLx, dtype: int64
```

```
In [6]: # Relabel the gender Status variable
da["RIAGENDRx"] = da.RIAGENDR.replace({1: "Male", 2: "Female"})
# Check the class distribution of RIAGENDR using groupby and size
print(da.groupby('RIAGENDRx').size())
```

```
RIAGENDRx
Female    2976
Male      2759
dtype: int64
```

**Q1a.** Briefly comment on some of the differences that you observe between the distribution of marital status between women and men, for people of all ages.

```
In [7]: dx = da.groupby(["DMDMARTLx"])["RIAGENDRx"].value_counts().unstack()
        dx = dx.apply(lambda x: x/x.sum(), axis = 0)
        print(dx.to_string(float_format = "%.4f"))
```

```
RIAGENDRx      Female    Male
DMDMARTLx
Divorced        0.1176  0.0830
Living with Partner 0.0880  0.0960
Married         0.4378  0.5353
Missing         0.0423  0.0489
Never Married   0.1747  0.1754
Refused         0.0003  0.0004
Separated       0.0397  0.0246
Widowed         0.0995  0.0362
```

**Q1b.** Briefly comment on the differences that you observe between the distribution of marital status states for women between the overall population, and for women between the ages of 30 and 40.

```
In [8]: # between the distribution of marital status status for women between the overall popu
da["agegap"] = pd.cut(da.RIDAGEYR, [30, 40])
da[(da.RIAGENDRx == "Female") & (da.agegap == pd.Interval(30, 40))].DMDMARTLx.value_counts()
```

```
Out[8]: Married          0.044987
        Never Married    0.016914
        Living with Partner 0.009939
        Divorced         0.007498
        Separated        0.002964
        Widowed          0.000349
        Name: DMDMARTLx, dtype: float64
```

```
In [9]: # between the distribution of marital status states for women between the ages of 30 and 40
da["agegap"] = pd.cut(da.RIDAGEYR,[30, 40])
da[(da.RIAGENDRx == "Female") & (da.agegap == pd.Interval(30, 40))].DMDMARTLx.value_counts()
```

```
Out[9]: Married          0.044987
        Never Married    0.016914
        Living with Partner 0.009939
        Divorced         0.007498
        Separated        0.002964
        Widowed          0.000349
        Name: DMDMARTLx, dtype: float64
```

**Q1c.** Repeat part b for the men.

```
In [10]: # between the distribution of marital status states for men between the overall population
da[(da.RIAGENDRx == "Male")].DMDMARTLx.value_counts()/da["DMDMARTLx"].shape[0]
```

```
Out[10]: Married          0.257541
          Never Married    0.084394
          Living with Partner 0.046207
          Divorced         0.039930
          Missing          0.023540
          Widowed          0.017437
          Separated        0.011857
          Refused          0.000174
          Name: DMDMARTLx, dtype: float64
```

```
In [11]: # between the distribution of marital status states for men between the ages of 30 and 40
da["agegap"] = pd.cut(da.RIDAGEYR,[30, 40])
da[(da.RIAGENDRx == "Male") & (da.agegap == pd.Interval(30, 40))].DMDMARTLx.value_counts()
```

```
Out[11]: Married          0.044987
          Never Married    0.015519
          Living with Partner 0.012554
          Divorced         0.004185
          Separated        0.002092
          Widowed          0.000349
          Refused          0.000174
          Name: DMDMARTLx, dtype: float64
```

## 1.2 Question 2

Restricting to the female population, stratify the subjects into age bands no wider than ten years, and construct the distribution of marital status within each age band. Within each age band, present the distribution in terms of proportions that must sum to 1.

In [16]: *# insert your code here*

```
x = da[da.RIAGENDRx == "Female"]
x["agegap2"] = pd.cut(da.RIDAGEYR, [10, 20, 30, 40, 50, 60, 70, 80])
dx = x.groupby(["agegap2"])["DMDMARTLx"].value_counts().unstack()
dx = dx.apply(lambda y: y/y.sum(), axis = 0)
print(dx.to_string(float_format = "%.2f"))
```

DMDMARTLx	Divorced	Living with Partner	Married	Missing	Never Married	Refused	Separated
agegap2							
(10, 20]	NaN	0.03	0.00	1.00	0.06	NaN	NaN
(20, 30]	0.03	0.40	0.12	NaN	0.44	NaN	0.09
(30, 40]	0.12	0.22	0.20	NaN	0.19	NaN	0.14
(40, 50]	0.20	0.14	0.22	NaN	0.12	NaN	0.28
(50, 60]	0.24	0.12	0.20	NaN	0.08	1.00	0.23
(60, 70]	0.24	0.07	0.16	NaN	0.07	NaN	0.19
(70, 80]	0.17	0.01	0.10	NaN	0.04	NaN	0.07

In [17]: `da["agegap"] = pd.cut(da.RIDAGEYR, [10, 20, 30, 40, 50, 60, 70, 80])`

`(da[da["RIAGENDRx"] == "Female"].groupby(["agegap", "DMDMARTLx"]).size() / da[da["RIAGENDRx"] == "Female"].size())`

Out[17]:

DMDMARTLx	Divorced	Living with Partner	Married	Missing	Never Married	\
agegap						
(10, 20]	NaN	0.048485	0.006061	0.763636	0.181818	
(20, 30]	0.021401	0.206226	0.305447	NaN	0.445525	
(30, 40]	0.090717	0.120253	0.544304	NaN	0.204641	
(40, 50]	0.137450	0.073705	0.573705	NaN	0.125498	
(50, 60]	0.176596	0.068085	0.546809	NaN	0.089362	
(60, 70]	0.192744	0.043084	0.480726	NaN	0.086168	
(70, 80]	0.143902	0.007317	0.317073	NaN	0.051220	

DMDMARTLx	Refused	Separated	Widowed
agegap			
(10, 20]	NaN	NaN	NaN
(20, 30]	NaN	0.021401	NaN
(30, 40]	NaN	0.035865	0.004219
(40, 50]	NaN	0.065737	0.023904
(50, 60]	0.002128	0.057447	0.059574
(60, 70]	NaN	0.049887	0.147392
(70, 80]	NaN	0.019512	0.460976

**Q2a.** Comment on the trends that you see in this series of marginal distributions.

**Q2b.** Repeat the construction for males.

In [18]: # insert your code here

```
x = da[da.RIAGENDRx == "Male"]
x["agegap2"] = pd.cut(da.RIDAGEYR, [10, 20, 30, 40, 50, 60, 70, 80])
dx = x.groupby(["agegap2"])["DMDMARTLx"].value_counts().unstack()
dx = dx.apply(lambda y: y/y.sum(), axis = 0)
print(dx.to_string(float_format = "%.2f"))
```

DMDMARTLx	Divorced	Living with Partner	Married	Missing	Never Married	Refused	Separated
agegap2							
(10, 20]	NaN	0.01	0.00	1.00	0.07	NaN	NaN
(20, 30]	0.01	0.35	0.07	NaN	0.47	NaN	0.10
(30, 40]	0.10	0.27	0.17	NaN	0.18	1.00	0.18
(40, 50]	0.15	0.12	0.19	NaN	0.08	NaN	0.16
(50, 60]	0.25	0.13	0.20	NaN	0.10	NaN	0.15
(60, 70]	0.24	0.08	0.20	NaN	0.08	NaN	0.21
(70, 80]	0.25	0.03	0.17	NaN	0.02	NaN	0.21

In [19]: da["agegap"] = pd.cut(da.RIDAGEYR, [10, 20, 30, 40, 50, 60, 70, 80])  
 (da[da["RIAGENDRx"] == "Male"].groupby(["agegap", "DMDMARTLx"]).size() / da[da["RIAGENDRx"] == "Male"].size())

DMDMARTLx	Divorced	Living with Partner	Married	Missing	Never Married	\
agegap						
(10, 20]	NaN	0.018182	0.006061	0.818182	0.218182	
(20, 30]	0.003891	0.178988	0.200389	NaN	0.439689	
(30, 40]	0.050633	0.151899	0.544304	NaN	0.187764	
(40, 50]	0.067729	0.065737	0.561753	NaN	0.077689	
(50, 60]	0.121277	0.072340	0.629787	NaN	0.100000	
(60, 70]	0.124717	0.049887	0.659864	NaN	0.086168	
(70, 80]	0.139024	0.021951	0.600000	NaN	0.021951	


DMDMARTLx	Refused	Separated	Widowed
agegap			
(10, 20]	NaN	NaN	NaN
(20, 30]	NaN	0.013619	0.003891
(30, 40]	0.00211	0.025316	0.004219
(40, 50]	NaN	0.021912	0.003984
(50, 60]	NaN	0.021277	0.021277
(60, 70]	NaN	0.031746	0.038549
(70, 80]	NaN	0.034146	0.163415

Q2c. Comment on any notable differences that you see when comparing these results for females and for males.

### 1.3 Question 3


Construct a histogram of the distribution of heights using the BMXHT variable in the NHANES sample.

```
In [23]: # insert your code here
sns.distplot(da.BMXHT.dropna())
plt.show()
```



**Q3a.** Use the `bins` argument to `distplot` to produce histograms with different numbers of bins. Assess whether the default value for this argument gives a meaningful result, and comment on what happens as the number of bins grows excessively large or excessively small.


```
In [24]: sns.distplot(da.BMXHT.dropna(), bins = 50)
plt.show()
```

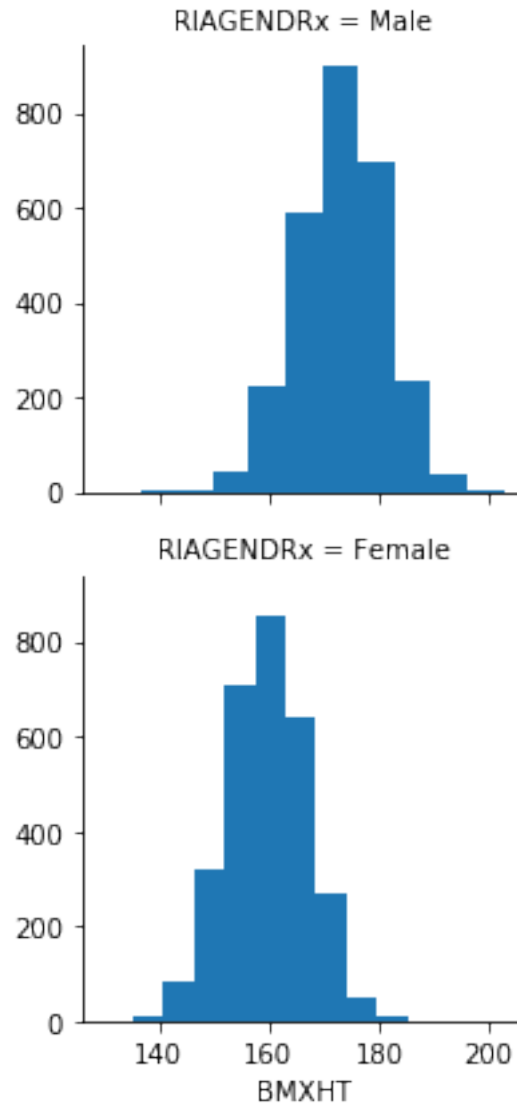


**Q3b.** Make separate histograms for the heights of women and men, then make a side-by-side boxplot showing the heights of women and men.

```
In [25]: # insert your code here
sns.boxplot(x= da["BMXHT"], y = da["RIAGENDRx"])
g = sns.FacetGrid(da, row = "RIAGENDRx")
g = g.map(plt.hist, "BMXHT")
plt.show()
```





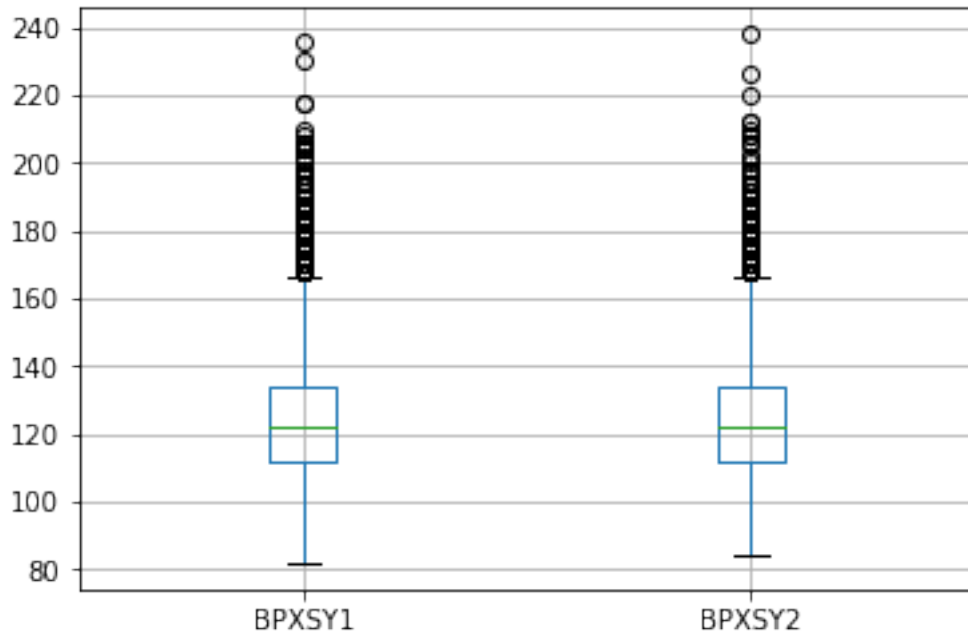


**Q3c.** Comment on what features, if any are not represented clearly in the boxplots, and what features, if any, are easier to see in the boxplots than in the histograms.

#### 1.4 Question 4

Make a boxplot showing the distribution of within-subject differences between the first and second systolic blood pressure measurements (BPXSY1 and BPXSY2).

```
In [26]: # insert your code here
x = da[["BPXSY1", "BPXSY2"]]
x.boxplot()
plt.show()
```



**Q4a.** What proportion of the subjects have a lower SBP on the second reading compared to the first?

In [ ]: *# insert your code here*

**Q4b.** Make side-by-side boxplots of the two systolic blood pressure variables.

In [4]: *# insert your code here*

**Q4c.** Comment on the variation within either the first or second systolic blood pressure measurements, and the variation in the within-subject differences between the first and second systolic blood pressure measurements.

## 1.5 Question 5

Construct a frequency table of household sizes for people within each educational attainment category (the relevant variable is `DMDEDUC2`). Convert the frequencies to proportions.

```
In [28]: # insert your code here
dx = da.groupby(["DMDEDUC2"])["DMDHHSIZ"].value_counts().unstack()
dx = dx.apply(lambda x: x/x.sum(), axis = 1)
print(dx.to_string(float_format="%.2f"))
```

DMDHHSIZ	1	2	3	4	5	6	7
DMDEDUC2							
1.0	0.11	0.22	0.15	0.13	0.15	0.11	0.13
2.0	0.12	0.22	0.16	0.15	0.15	0.11	0.09

3.0	0.15	0.27	0.17	0.16	0.11	0.07	0.07
4.0	0.15	0.27	0.19	0.17	0.12	0.05	0.05
5.0	0.14	0.35	0.19	0.17	0.10	0.03	0.03
9.0	NaN	0.67	NaN	NaN	0.33	NaN	NaN

**Q5a.** Comment on any major differences among the distributions.

**Q5b.** Restrict the sample to people between 30 and 40 years of age. Then calculate the median household size for women and men within each level of educational attainment.

In [29]: *# insert your code here*

```
da[(da.RIDAGEYR >= 30) & (da.RIDAGEYR <= 40)].groupby(["DMDEDUC2", "RIAGENDR"])["DMDHHSIZ"]
```

Out [29]:

DMDEDUC2	RIAGENDR	DMDHHSIZ
1.0	1	5.0
	2	5.0
2.0	1	4.5
	2	5.0
3.0	1	4.0
	2	5.0
4.0	1	4.0
	2	4.0
5.0	1	3.0
	2	3.0

Name: DMDHHSIZ, dtype: float64

## 1.6 Question 6

The participants can be clustered into “maked variance units” (MVU) based on every combination of the variables `SDMVSTRA` and `SDMVPSU`. Calculate the mean age (`RIDAGEYR`), height (`BMXHT`), and BMI (`BMXBMI`) for each gender (`RIAGENDR`), within each MVU, and report the ratio between the largest and smallest mean (e.g. for height) across the MVUs.

In [36]: *# insert your code here*

```
(
    (
        da.groupby(['SDMVSTRA', 'SDMVPSU', 'RIAGENDR'])
        [['RIDAGEYR', 'BMXHT', 'BMXBMI']]
        .mean()
    )
).unstack()
```

Out [36]:

		RIDAGEYR		BMXHT		BMXBMI	\
RIAGENDR		1	2	1	2	1	
SDMVSTRA	SDMVPSU						
119	1	47.861111	47.663265	172.741667	159.570408	26.958333	
	2	54.363636	52.987952	172.906818	159.244578	27.160465	
120	1	43.130000	43.636364	169.537755	155.402041	30.939175	

	2	45.219178	43.736111	173.075342	159.218056	27.727397
121	1	46.750000	44.397959	172.177885	158.871579	29.416505
	2	42.063158	44.376344	174.764516	160.229032	26.273118
122	1	44.653061	42.897436	173.998969	161.315385	28.528866
	2	44.320000	47.333333	170.332323	157.231111	25.744444
123	1	47.829787	44.841121	174.315217	162.059615	29.231522
	2	52.126582	46.457447	174.454430	160.476596	28.811392
124	1	50.750000	51.664000	172.109009	158.788710	28.614414
	2	48.245614	42.541667	174.291228	162.853521	27.714035
125	1	55.165289	50.900901	173.631092	160.762385	29.727731
	2	49.705882	51.660000	174.456863	160.021429	29.143564
126	1	48.416667	46.229167	175.149398	160.387500	29.033333
	2	48.666667	47.205882	174.713043	160.892000	29.039130
127	1	53.137931	49.694444	171.545349	157.422430	31.062353
	2	54.070588	51.486239	173.366667	159.022936	30.557831
128	1	53.673267	55.638462	169.325000	156.339062	31.749000
	2	45.822785	45.589744	172.400000	160.437179	26.835443
129	1	43.922222	45.329787	171.094318	156.900000	26.493182
	2	45.775510	43.500000	173.138298	161.034259	28.961702
130	1	50.516854	47.810526	176.974157	161.977895	30.337079
	2	50.535354	50.833333	175.061224	160.060577	29.237755
131	1	53.140187	54.893617	175.610476	161.989362	28.259615
	2	46.778846	45.000000	175.091346	161.673810	30.077885
132	1	42.380435	43.210526	172.534066	161.508421	28.546154
	2	49.038760	51.700000	172.809524	159.138281	28.966667
133	1	44.054795	45.105882	171.509722	158.295122	27.495833
	2	47.489796	47.063158	171.179167	158.627368	27.966667

RIAGENDR	2	
SDMVSTRA	SDMVPSU	
119	1	30.052041
	2	27.849398
120	1	32.419388
	2	27.400000
121	1	30.856842
	2	26.470968
122	1	29.447436
	2	26.611111
123	1	29.905769
	2	30.641489
124	1	29.533065
	2	28.640845
125	1	30.385321
	2	28.564286
126	1	31.262500
	2	29.612121
127	1	32.189720

	2	30.770642
128	1	32.303125
	2	27.491026
129	1	29.019149
	2	29.429630
130	1	30.700000
	2	31.490385
131	1	30.061702
	2	32.984127
132	1	29.848421
	2	30.540625
133	1	27.959259
	2	29.000000

```
In [38]: (
    (
        da.groupby(['SDMVSTRA', 'SDMVPSU', 'RIAGENDR'])
        [['RIDAGEYR', 'BMXHT', 'BMXBMI']]
        .max()
    )
    /
    (
        da.groupby(['SDMVSTRA', 'SDMVPSU', 'RIAGENDR'])
        [['RIDAGEYR', 'BMXHT', 'BMXBMI']]
        .min()
    )
).unstack()
```

```
Out [38]:
```

		RIDAGEYR		BMXHT		BMXBMI	
		1	2	1	2	1	2
RIAGENDR	SDMVSTRA	SDMVPSU					
119	1	4.444444	4.444444	1.231270	1.221838	2.994413	3.256684
	2	4.000000	4.444444	1.269385	1.267041	2.224390	4.045161
120	1	4.444444	4.444444	1.297715	1.254360	3.042105	3.666667
	2	4.444444	4.444444	1.271065	1.296011	2.875000	2.443182
121	1	4.444444	4.444444	1.252324	1.308458	3.104938	2.538462
	2	4.444444	4.444444	1.296036	1.249319	2.542373	3.349112
122	1	4.444444	4.444444	1.244646	1.232877	3.284916	3.066298
	2	4.444444	4.444444	1.272849	1.214035	1.902703	2.730539
123	1	4.444444	4.444444	1.210127	1.262143	3.039548	2.483333
	2	4.210526	4.444444	1.290116	1.237474	3.275000	2.947917
124	1	4.210526	4.444444	1.214700	1.259393	2.212871	3.656627
	2	4.210526	4.444444	1.204545	1.216480	2.436464	3.530387
125	1	4.444444	4.444444	1.266364	1.230070	3.271605	2.948571
	2	4.444444	4.444444	1.229072	1.248281	2.420765	3.437126
126	1	4.210526	4.210526	1.224347	1.223684	2.769231	3.395210
	2	4.444444	4.444444	1.244715	1.217753	2.355670	3.226994
127	1	4.444444	4.444444	1.219481	1.226573	3.220930	3.195531

	2	4.444444	4.444444	1.223082	1.267626	2.704142	3.184971
128	1	4.444444	4.000000	1.359029	1.277698	2.577114	3.160428
	2	4.444444	4.444444	1.271883	1.223135	2.734940	3.434483
129	1	4.444444	4.444444	1.209941	1.350810	2.124324	2.806630
	2	4.444444	4.444444	1.250484	1.245775	3.039326	3.413408
130	1	4.444444	4.444444	1.224691	1.227972	3.457831	3.261538
	2	4.444444	4.000000	1.255995	1.288210	2.994475	3.875862
131	1	4.210526	4.444444	1.203232	1.220938	3.357616	3.444444
	2	4.444444	4.444444	1.204759	1.203528	2.713568	3.822485
132	1	4.111111	4.210526	1.268568	1.198903	2.432161	3.121212
	2	4.444444	4.444444	1.292895	1.275562	3.512195	4.078788
133	1	4.444444	4.444444	1.265293	1.273934	2.502762	2.969512
	2	4.444444	4.444444	1.372161	1.202364	3.222222	3.706897

**Q6a.** Comment on the extent to which mean age, height, and BMI vary among the MVUs.

**Q6b.** Calculate the inter-quartile range (IQR) for age, height, and BMI for each gender and each MVU. Report the ratio between the largest and smallest IQR across the MVUs.

In [39]: # insert your code here

```
(
    (
        da.groupby(['SDMVSTRA', 'SDMVPSU', 'RIAGENDR'])
        [['RIDAGEYR', 'BMXHT', 'BMXBMI']]
        .quantile(0.75)
    )
    -
    (
        da.groupby(['SDMVSTRA', 'SDMVPSU', 'RIAGENDR'])
        [['RIDAGEYR', 'BMXHT', 'BMXBMI']]
        .quantile(0.25)
    )
).unstack()
```

Out [39]:

			RIDAGEYR		BMXHT		BMXBMI	
			1	2	1	2	1	2
	SDMVSTRA	SDMVPSU						
119	1		29.75	31.25	9.000	9.325	5.350	9.750
	2		29.00	33.50	11.225	9.950	5.300	9.350
120	1		23.75	26.50	12.125	8.750	9.400	8.775
	2		26.00	25.75	10.500	10.550	7.100	7.750
121	1		34.50	26.25	10.725	9.150	7.500	9.000
	2		25.50	26.00	8.600	9.600	5.700	8.100
122	1		29.50	24.00	9.400	10.400	7.700	9.875
	2		30.00	25.00	10.150	7.575	4.100	8.475
123	1		28.25	30.50	9.350	9.675	8.050	10.450
	2		31.50	34.50	9.900	11.200	8.100	9.975
124	1		32.00	27.00	9.800	8.375	6.100	8.950
	2		31.00	23.50	11.600	8.650	8.700	9.000

125	1	29.00	31.00	10.350	9.100	8.300	8.000
	2	33.50	32.25	7.925	10.675	7.900	10.325
126	1	36.25	30.25	10.450	8.500	8.000	10.675
	2	34.00	31.75	8.125	12.025	6.850	10.350
127	1	30.00	27.25	9.025	7.700	8.200	11.750
	2	28.00	30.00	10.750	11.600	5.950	9.200
128	1	33.00	28.00	9.950	9.125	6.675	8.500
	2	25.50	22.00	9.850	10.650	5.800	9.375
129	1	20.75	24.75	12.300	10.375	6.025	9.500
	2	30.75	26.25	10.700	8.900	5.800	9.725
130	1	36.00	35.50	9.900	8.650	6.700	11.200
	2	28.50	30.25	8.625	10.225	8.375	8.050
131	1	36.00	35.75	10.500	10.025	7.525	11.075
	2	28.00	24.00	7.750	7.575	7.850	10.625
132	1	21.25	30.00	10.600	10.950	6.600	10.700
	2	38.00	33.00	10.550	10.100	9.600	11.750
133	1	33.00	34.00	8.925	10.300	6.425	8.300
	2	32.25	28.50	8.850	9.550	5.900	9.650

**Q6c.** Comment on the extent to which the IQR for age, height, and BMI vary among the MVUs.